

**Aim:** To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

### **Theory:**

Container-based microservices architectures have revolutionized how development and operations teams test and deploy modern software. Containers allow companies to scale and deploy applications more efficiently, but they also introduce new challenges, adding complexity by creating a whole new infrastructure ecosystem.

Today, both large and small software companies are deploying thousands of container instances daily. Managing this level of complexity at scale requires advanced tools. Enter Kubernetes.

Originally developed by Google, Kubernetes is an open-source container orchestration platform designed to automate the deployment, scaling, and management of containerized applications. Kubernetes has quickly become the de facto standard for container orchestration and is the flagship project of the Cloud Native Computing Foundation (CNCF), supported by major players like Google, AWS, Microsoft, IBM, Intel, Cisco, and Red Hat.

Kubernetes simplifies the deployment and operation of applications in a microservice architecture by providing an abstraction layer over a group of hosts. This allows development teams to deploy their applications while Kubernetes takes care of key tasks, including:

- Managing resource consumption by applications or teams
- Distributing application load evenly across the infrastructure
- Automatically load balancing requests across multiple instances of an application
- Monitoring resource usage to prevent applications from exceeding resource limits and automatically restarting them if needed
- Moving application instances between hosts when resources are low or if a host fails
- Automatically utilizing additional resources when new hosts are added to the cluster
- Facilitating canary deployments and rollbacks with ease

### **Necessary Requirements:**

• **EC2 Instance:** The experiment required launching a t2.medium EC2 instance with 2 CPUs, as Kubernetes demands sufficient resources for effective functioning.

#### • **Minimum Requirements:**

- **Instance Type:** t2.medium
- **CPUs:** 2
- **Memory:** Adequate for container orchestration.

This ensured that the Kubernetes cluster had the necessary resources to function smoothly

#### **Note:**

AWS Personal Account is preferred but we can also perform it on AWS Academy (adding some ignores in the command if any error occurs in below as the below experiment is performed on Personal Account). If You are using AWS Academy Account Errors you will face in kubeadm init command so you have to add some ignores with this command.

### Prerequisites :

**Create 2 Security Groups for Master and Nodes and add the following rules inbound rules in those Groups.**

**Master:**

Inbound rules <a href="#">info</a>						
Security group rule ID	Type <a href="#">info</a>	Protocol <a href="#">info</a>	Port range <a href="#">info</a>	Source <a href="#">info</a>	Description - optional <a href="#">info</a>	
sgr-0c17c1a22a7c7b3e5	HTTP	TCP	80	Custom	<input type="text" value="Q"/>	<input type="button" value="Delete"/>
					<input type="text" value="0.0.0.0"/> <input type="button" value="X"/>	
sgr-0d3f86194443b29f1	All traffic	All	All	Custom	<input type="text" value="Q"/>	<input type="button" value="Delete"/>
					<input type="text" value="0.0.0.0"/> <input type="button" value="X"/>	
sgr-010d128b1484f322	Custom TCP	TCP	6443	Custom	<input type="text" value="Q"/>	<input type="button" value="Delete"/>
					<input type="text" value="0.0.0.0"/> <input type="button" value="X"/>	
sgr-05bb413f0626b9c3b	Custom TCP	TCP	10251	Custom	<input type="text" value="Q"/>	<input type="button" value="Delete"/>
					<input type="text" value="0.0.0.0"/> <input type="button" value="X"/>	
sgr-04bd098c8f409420d	Custom TCP	TCP	10250	Custom	<input type="text" value="Q"/>	<input type="button" value="Delete"/>
					<input type="text" value="0.0.0.0"/> <input type="button" value="X"/>	
sgr-01438a40425cf867c	All TCP	TCP	0 - 65535	Custom	<input type="text" value="Q"/>	<input type="button" value="Delete"/>
					<input type="text" value="0.0.0.0"/> <input type="button" value="X"/>	
sgr-05dc20e8c2b541402	Custom TCP	TCP	10252	Custom	<input type="text" value="Q"/>	<input type="button" value="Delete"/>
					<input type="text" value="0.0.0.0"/> <input type="button" value="X"/>	
sgr-08d45afafe6c06c26	SSH	TCP	22	Custom	<input type="text" value="Q"/>	<input type="button" value="Delete"/>
					<input type="text" value="0.0.0.0"/> <input type="button" value="X"/>	

**Node :**

**Edit inbound rules** [Info](#)

Inbound rules control the incoming traffic that's allowed to reach the instance.

### Inbound rules [Info](#)

Security group rule ID	Type <a href="#">Info</a>	Protocol <a href="#">Info</a>	Port range <a href="#">Info</a>	Source <a href="#">Info</a>	Description - optional <a href="#">Info</a>	
sgr-00d83454961e5d3e9	All traffic	All	All	Custom	<input type="text"/>	Delete
				<input type="text"/> 0.0.0.0/0		
sgr-0402e9e84cd3dea45	SSH	TCP	22	Custom	<input type="text"/>	Delete
				<input type="text"/> 0.0.0.0/0		
sgr-05770af1e4c5669f7	Custom TCP	TCP	10250	Custom	<input type="text"/>	Delete
				<input type="text"/> 0.0.0.0/0		
sgr-0b3fb7516970bc90	All TCP	TCP	0 - 65535	Custom	<input type="text"/>	Delete
				<input type="text"/> 0.0.0.0/0		
sgr-07384bc31bec899e9	Custom TCP	TCP	30000 - 32767	Custom	<input type="text"/>	Delete
				<input type="text"/> 0.0.0.0/0		
sgr-05188e46d7e21828d	HTTP	TCP	80	Custom	<input type="text"/>	Delete
				<input type="text"/> 0.0.0.0/0		

Add rule

Name: Bhushan Mukund Kor

Academic Year: 2024-2025

Division: D15C

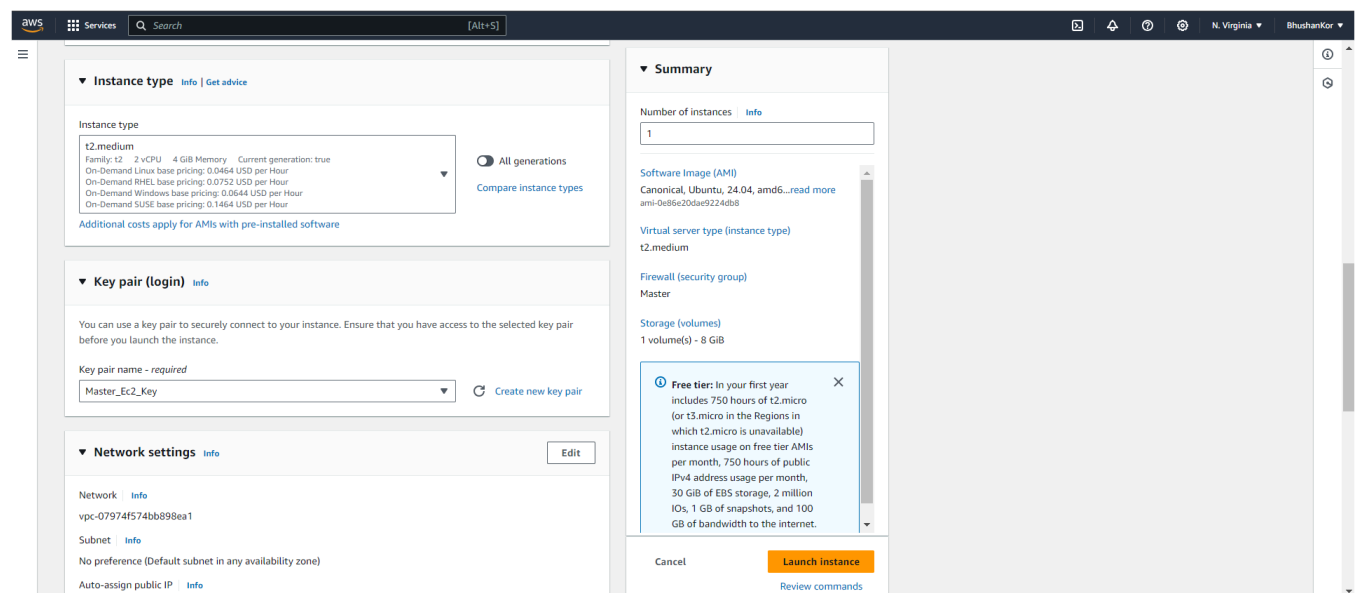
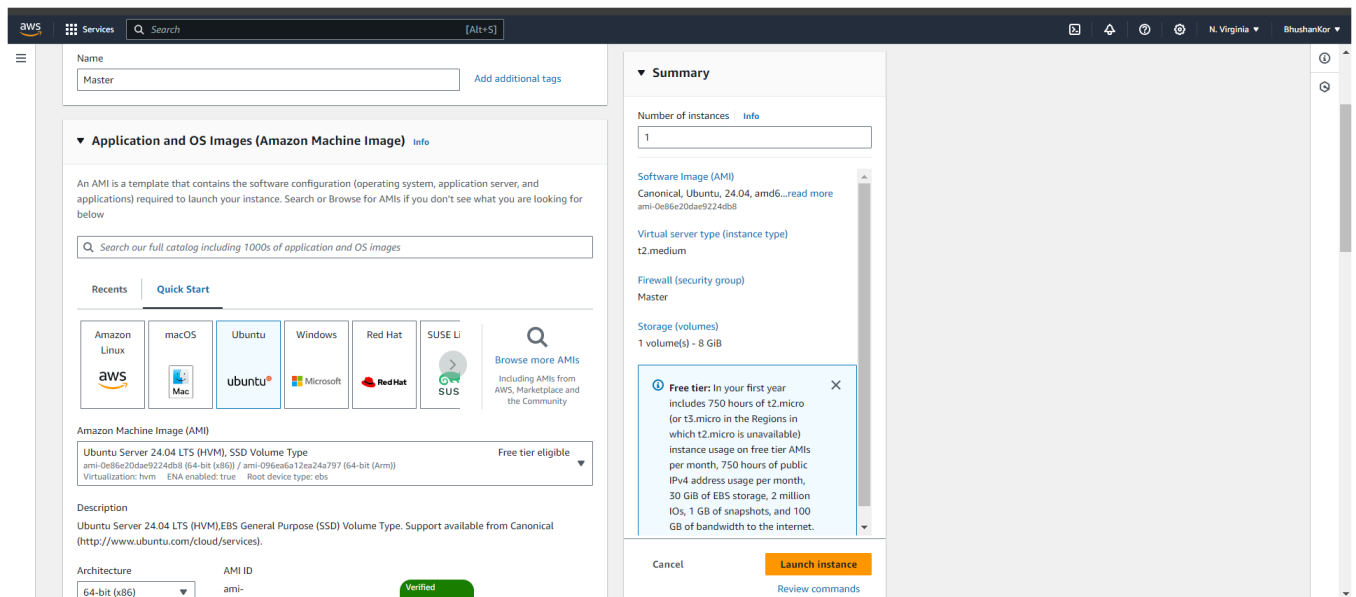
Roll No: 28

**Step 1:** Log in to your AWS Academy/personal account and launch 3 new Ec2 Instances. Select Ubuntu as AMI and **t2.medium** as Instance Type and create a key of type RSA with .pem extension and move the downloaded key to the new folder. We can use 3 Different keys or 1 common key also.

**Note:** A minimum of 2 CPUs are required so Please select t2.medium and do not forget to stop the instance after the experiment because it is not available in the free tier.

**Also Select Security groups from existing.**

**Master:**



Name: Bhushan Mukund Kor

Academic Year: 2024-2025

Division: D15C

Roll No: 28

**Network settings** [Info](#) [Edit](#)

**Network** [Info](#)  
vpc-07974f574bb898ea1

**Subnet** [Info](#)  
No preference (Default subnet in any availability zone)

**Auto-assign public IP** [Info](#)  
Enable

**Additional charges apply** when outside of free tier allowance

**Firewall (security group)** [Info](#)  
A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☐ Create security group ☒ Select existing security group

**Common security groups** [Info](#)

Select security groups

☐ Node VPC: vpc-07974f574bb898ea1 sg-0990b1794d851ae05

☒ Master VPC: vpc-07974f574bb898ea1 sg-0ab4c57c9d569b1c8

☐ default VPC: vpc-07974f574bb898ea1 sg-086e3eb335693fc6a

[Compare security group rules](#)

**Storage (volumes)**  
1 volume(s) - 8 GiB

**Free tier:** In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GiB of snapshots, and 100 GB of bandwidth to the internet.

[Launch instance](#) [Review commands](#)

Do Same for 2 Nodes and use security groups of Node for that.

**Step 2:** After creating the instances click on Connect & connect all 3 instances and navigate to SSH Client.

Instances (3) <a href="#">Info</a>										
<input type="text" value="Find Instance by attribute or tag (case-sensitive)"/>					Running					
	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP
<input type="checkbox"/>	Master	i-0cfb6b3b53bc03ab1	Running	t2.medium	2/2 checks passed	<a href="#">View alarms</a>	us-east-1d	ec2-52-90-3-215.comp...	52.90.3.215	-
<input type="checkbox"/>	Node 1	i-0b64c605d31b0bd44	Running	t2.medium	Initializing	<a href="#">View alarms</a>	us-east-1d	ec2-3-80-56-65.comput...	3.80.56.65	-
<input type="checkbox"/>	Node 2	i-0af54010ae84808d2	Running	t2.medium	Initializing	<a href="#">View alarms</a>	us-east-1d	ec2-34-224-169-38.co...	34.224.169.38	-

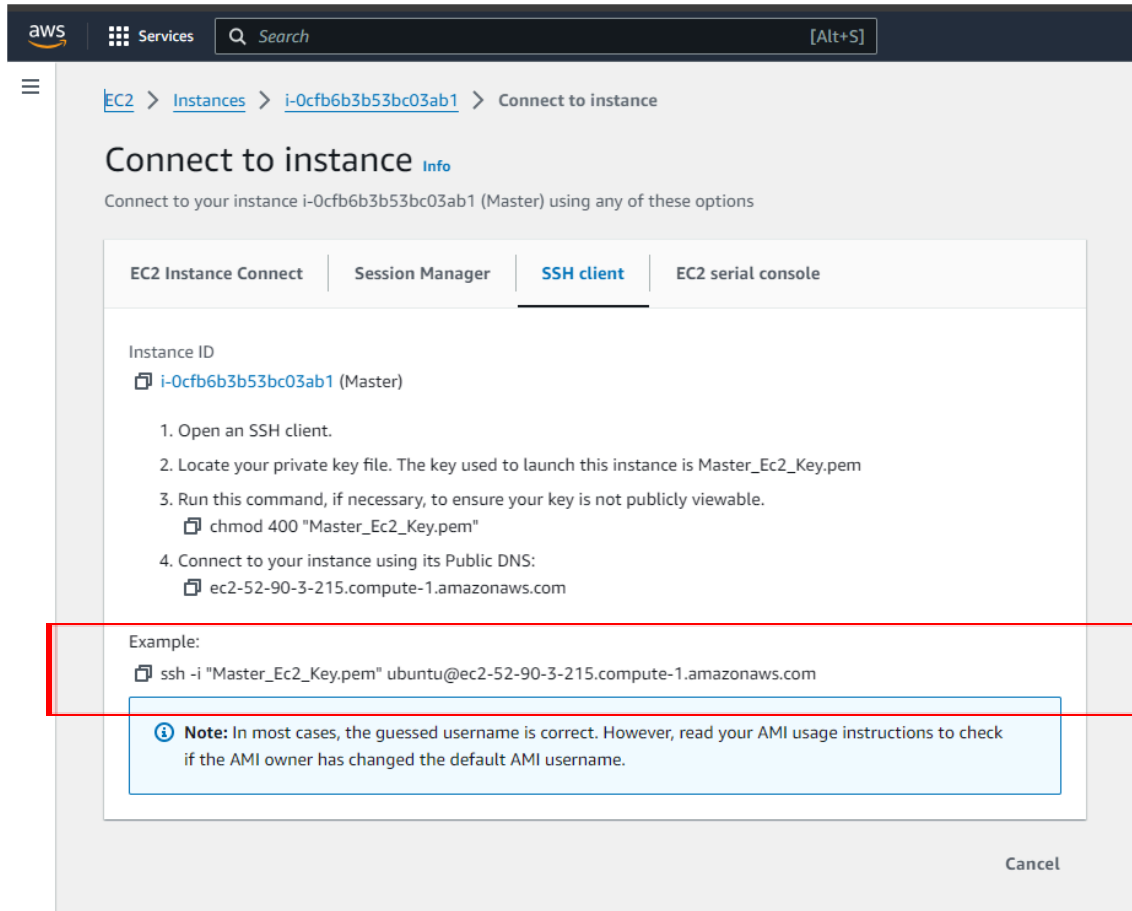
(Downloaded Key

Bhushan - Personal > Desktop > New folder (4)						Sea
Sort View ...						
Name	Status	Date modified	Type	Size		
Master_Ec2_Key.pem		9/14/2024 4:32 PM	PEM File	2 KB		

)

**Step 3:** Now open the folder in the terminal 3 times for Master, Node1 & Node 2 where our .pem key is stored and paste the Example command (starting with `ssh -i .....`) in the terminal. (`ssh -i "Master_Ec2_Key.pem" ubuntu@ec2-54-196-129-215.compute-1.amazonaws.com`)

Master:



aws Services Search [Alt+S]

EC2 > Instances > i-0cfb6b3b53bc03ab1 > Connect to instance

### Connect to instance Info

Connect to your instance i-0cfb6b3b53bc03ab1 (Master) using any of these options

EC2 Instance Connect | Session Manager | **SSH client** | EC2 serial console

Instance ID  
i-0cfb6b3b53bc03ab1 (Master)

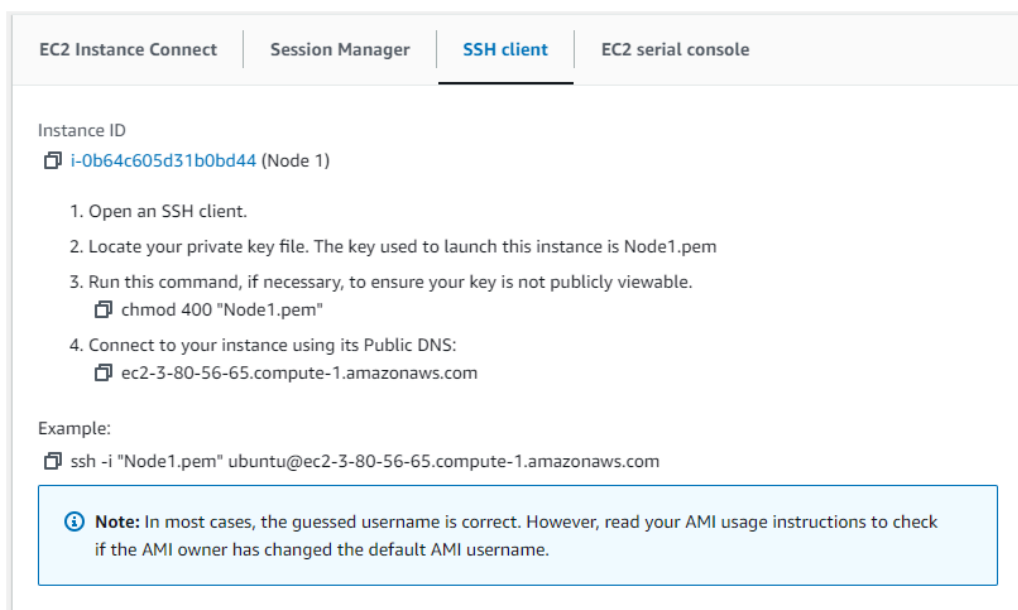
1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is Master\_Ec2\_Key.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.  
`chmod 400 "Master_Ec2_Key.pem"`
4. Connect to your instance using its Public DNS:  
`ec2-52-90-3-215.compute-1.amazonaws.com`

Example:  
`ssh -i "Master_Ec2_Key.pem" ubuntu@ec2-52-90-3-215.compute-1.amazonaws.com`

**Note:** In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Cancel

Node 1:



EC2 Instance Connect | Session Manager | **SSH client** | EC2 serial console

Instance ID  
i-0b64c605d31b0bd44 (Node 1)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is Node1.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.  
`chmod 400 "Node1.pem"`
4. Connect to your instance using its Public DNS:  
`ec2-3-80-56-65.compute-1.amazonaws.com`

Example:  
`ssh -i "Node1.pem" ubuntu@ec2-3-80-56-65.compute-1.amazonaws.com`

**Note:** In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Node 2:

EC2 Instance Connect

Session Manager

SSH client

EC2 serial console

Instance ID

i-0af54010ae84808d2 (Node 2)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is Node1.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.

chmod 400 "Node1.pem"
4. Connect to your instance using its Public DNS:

ec2-34-224-169-38.compute-1.amazonaws.com

Example:

ssh -i "Node1.pem" ubuntu@ec2-34-224-169-38.compute-1.amazonaws.com

**Note:** In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Here I have use 2 keys 1 for master and 1 for 2 node so I have to run open 3 terminals. In master key folder 1 terminal and 2 terminals in node 1 key folder.

If you use 1 Key only, you can open 3 terminal in one folder only.

Successful Connection:

ubuntu@ip-172-31-27-176: ~

E.

This key is not known by any other names

Are you sure you want to continue connecting (yes/no/[fingerprint])? yes

Warning: Permanently added 'ec2-52-90-3-215.compute-1.amazonaws.com' (ED25519) to the list of known hosts.

Welcome to Ubuntu 24.04 LTS (GNU/Linux 6.8.0-1012-aws x86\_64)

\* Documentation: <https://help.ubuntu.com>

\* Management: <https://landscape.canonical.com>

\* Support: <https://ubuntu.com/pro>

System information as of Mon Sep 16 15:13:30 UTC 2024

System load: 0.08

Processes: 115

Usage of /: 22.9% of 6.71GB

Users logged in: 0

Memory usage: 5%

IPv4 address for enX0: 172.31.27.176

Swap usage: 0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates. See <https://ubuntu.com/esm> or run: sudo pro status

The list of available updates is more than a week old. To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software; the exact distribution terms for each program are described in the individual files in /usr/share/doc/\*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by applicable law.

To run a command as administrator (user "root"), use "sudo <command>". See "man sudo\_root" for details.

ubuntu@ip-172-31-27-176:~\$

ubuntu@ip-172-31-28-117: ~

Enable ESM Apps to receive additional future security updates. See <https://ubuntu.com/esm> or run: sudo pro status

The list of available updates is more than a week old. To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software; the exact distribution terms for each program are described in the individual files in /usr/share/doc/\*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by applicable law.

To run a command as administrator (user "root"), use "sudo <command>". See "man sudo\_root" for details.

ubuntu@ip-172-31-28-117:~\$

ubuntu@ip-172-31-18-135: ~

Enable ESM Apps to receive additional future security updates. See <https://ubuntu.com/esm> or run: sudo pro status

The list of available updates is more than a week old. To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software; the exact distribution terms for each program are described in the individual files in /usr/share/doc/\*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by applicable law.

To run a command as administrator (user "root"), use "sudo <command>". See "man sudo\_root" for details.

ubuntu@ip-172-31-18-135:~\$

**Step 4:** Run on Master, Node 1, and Node 2 the below commands to install and setup Docker in Master, Node1, and Node2.

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -  
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee  
/etc/apt/trusted.gpg.d/docker.gpg > /dev/null
```

```
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu  
$(lsb_release -cs) stable"
```

```
ubuntu@ip-172-31-27-176:~$  
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -  
  
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee /etc/apt/  
trusted.gpg.d/docker.gpg > /dev/null  
  
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/  
ubuntu $(lsb_release -cs) stable"  
Warning: apt-key is deprecated. Manage keyring files in trusted.gpg.d instea  
d (see apt-key(8)).  
OK  
Repository: 'deb [arch=amd64] https://download.docker.com/linux/ubuntu noble  
stable'  
Description:  
Archive for codename: noble components: stable  
More info: https://download.docker.com/linux/ubuntu  
Adding repository.  
Press [ENTER] to continue or Ctrl-c to cancel.  
Adding deb entry to /etc/apt/sources.list.d/archive_uri-https_download_docke  
r_com_linux_ubuntu-noble.list  
Adding disabled deb-src entry to /etc/apt/sources.list.d/archive_uri-https_d  
ownload_docker_com_linux_ubuntu-noble.list  
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease  
Get:2 https://download.docker.com/linux/ubuntu noble InRelease [48.8 kB]  
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease  
[126 kB]  
Get:4 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelea  
se [126 kB]  
Get:5 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 Pa  
ckages [15.0 MB]  
Get:6 https://download.docker.com/linux/ubuntu noble/stable amd64 Packages [ 13.8 kB]  
Get:7 http://security.ubuntu.com/ubuntu noble-security InRelease [126 kB]  
Get:8 http://security.ubuntu.com/ubuntu noble-security/main amd64 Packages [ 354 kB]  
Get:9 http://security.ubuntu.com/ubuntu noble-security/main Translation-en [ 79.4 kB]
```

```
Get:50 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports/universe amd64 c-n-f Metadata [116 B]
Get:51 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports/multiverse amd64 Components [212 B]
Get:52 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports/multiverse amd64 c-n-f Metadata [116 B]
Fetched 28.9 MB in 4s (6976 kB/s)
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: The key(s) in the keyring /etc/apt/trusted.gpg.d/docker.gpg are ignored as the file has an unsupported filetype.
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
ubuntu@ip-172-31-27-176:~$ |
```

**sudo apt-get update**

**sudo apt-get install -y docker-ce**

```
ubuntu@ip-172-31-27-176:~$ sudo apt-get update
sudo apt-get install -y docker-ce
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Hit:5 http://security.ubuntu.com/ubuntu noble-security InRelease
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: The key(s) in the keyring /etc/apt/trusted.gpg.d/docker.gpg are ignored as the file has an unsupported filetype.
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  containerd.io docker-buildx-plugin docker-ce-cli
  docker-ce-rootless-extras docker-compose-plugin libltdl7 libslirp0 pigz
  slirp4netns
Suggested packages:
  aufs-tools cgroupfs-mount | cgroup-lite
The following NEW packages will be installed:
  containerd.io docker-buildx-plugin docker-ce docker-ce-cli
  docker-ce-rootless-extras docker-compose-plugin libltdl7 libslirp0 pigz
  slirp4netns
0 upgraded, 10 newly installed, 0 to remove and 133 not upgraded.
Need to get 122 MB of archives.
After this operation, 440 MB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 pigz amd64 2.8-1 [65.6 kB]
```



```
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.  
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.  
Processing triggers for man-db (2.12.0-4build2) ...  
Processing triggers for libc-bin (2.39-0ubuntu8.2) ...  
Scanning processes...  
Scanning linux images...
```

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.  
ubuntu@ip-172-31-27-176:~\$ |

```
sudo mkdir -p /etc/docker  
cat <<EOF | sudo tee /etc/docker/daemon.json  
{  
  "exec-opts": ["native.cgroupdriver=systemd"]  
}  
EOF
```

```
ubuntu@ip-172-31-27-176:~$ sudo mkdir -p /etc/docker  
cat <<EOF | sudo tee /etc/docker/daemon.json  
{  
  "exec-opts": ["native.cgroupdriver=systemd"]  
}  
EOF  
{  
  "exec-opts": ["native.cgroupdriver=systemd"]  
}  
ubuntu@ip-172-31-27-176:~$
```

```
sudo systemctl enable docker  
sudo systemctl daemon-reload  
sudo systemctl restart docker
```

```
ubuntu@ip-172-31-27-176:~$ sudo systemctl enable docker  
sudo systemctl daemon-reload  
sudo systemctl restart docker  
Synchronizing state of docker.service with SysV service script with /usr/lib/systemd/systemd-sysv-install.  
Executing: /usr/lib/systemd/systemd-sysv-install enable docker  
ubuntu@ip-172-31-27-176:~$
```

**Step 5:** Run the below command to install Kubernetes.

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
```

```
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
```

```
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
```

```
ubuntu@ip-172-31-27-176:~$ curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
```

```
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
```

```
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /
ubuntu@ip-172-31-27-176:~$
```

**sudo apt-get update**

**sudo apt-get install -y kubelet kubeadm kubectl**

**sudo apt-mark hold kubelet kubeadm kubectl**

```
ubuntu@ip-172-31-27-176:~$ sudo apt-get update
sudo apt-get install -y kubelet kubeadm kubectl
sudo apt-mark hold kubelet kubeadm kubectl
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Hit:5 http://security.ubuntu.com/ubuntu noble-security InRelease
Get:6 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb InRelease [1186 B]
Get:7 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb Packages [4865 B]
Fetched 6051 B in 1s (11.1 kB/s)
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: The key(s) in the keyring /etc/apt/trusted.gpg.d/docker.gpg are ignored as the file has an unsupported filetype.
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
Reading package lists... Done
Building dependency tree... Done
```

```
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
kubelet set on hold.
kubeadm set on hold.
kubectl set on hold.
ubuntu@ip-172-31-27-176:~$ |
```

**sudo systemctl enable --now kubelet**

**sudo apt-get install -y containerd**

```
ubuntu@ip-172-31-27-176:~$ sudo systemctl enable --now kubelet
ubuntu@ip-172-31-27-176:~$ sudo apt-get install -y containerd
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras
  docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following additional packages will be installed:
  runc
The following packages will be REMOVED:
  containerd.io docker-ce
The following NEW packages will be installed:
  containerd runc
0 upgraded, 2 newly installed, 2 to remove and 133 not upgraded.
Need to get 47.2 MB of archives.
After this operation, 53.1 MB disk space will be freed.
```

```
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
```

**sudo mkdir -p /etc/containerd**

**sudo containerd config default | sudo tee /etc/containerd/config.toml**

```
ubuntu@ip-172-31-27-176:~$ sudo mkdir -p /etc/containerd
sudo containerd config default | sudo tee /etc/containerd/config.toml
disabled_plugins = []
imports = []
oom_score = 0
plugin_dir = ""
required_plugins = []
root = "/var/lib/containerd"
state = "/run/containerd"
temp = ""
version = 2

[cgroup]
  path = ""

[debug]
  address = ""
  format = ""
  gid = 0
  level = ""
  uid = 0

[grpc]
  address = "/run/containerd/containerd.sock"
  gid = 0
```

```
[timeouts]
  "io.containerd.timeout.bolt.open" = "0s"
  "io.containerd.timeout.metrics.shimstats" = "2s"
  "io.containerd.timeout.shim.cleanup" = "5s"
  "io.containerd.timeout.shim.load" = "5s"
  "io.containerd.timeout.shim.shutdown" = "3s"
  "io.containerd.timeout.task.state" = "2s"

[ttrpc]
  address = ""
  gid = 0
  uid = 0
```

**sudo systemctl restart containerd**

**sudo systemctl enable containerd**

**sudo systemctl status containerd**

```
ubuntu@ip-172-31-27-176:~$ sudo systemctl restart containerd
sudo systemctl enable containerd
sudo systemctl status containerd
● containerd.service - containerd container runtime
   Loaded: loaded (/usr/lib/systemd/system/containerd.service; enabled; p>
   Active: active (running) since Mon 2024-09-16 15:31:58 UTC; 210ms ago
     Docs: https://containerd.io
   Main PID: 4763 (containerd)
      Tasks: 7
   Memory: 13.9M (peak: 14.4M)
        CPU: 50ms
   CGroup: /system.slice/containerd.service
           └─4763 /usr/bin/containerd

Sep 16 15:31:58 ip-172-31-27-176 containerd[4763]: time="2024-09-16T15:31:5>
Sep 16 15:31:58 ip-172-31-27-176 containerd[4763]: time="2024-09-16T15:31:5>
Sep 16 15:31:58 ip-172-31-27-176 containerd[4763]: time="2024-09-16T15:31:5>
Sep 16 15:31:58 ip-172-31-27-176 containerd[4763]: time="2024-09-16T15:31:5>
Sep 16 15:31:58 ip-172-31-27-176 containerd[4763]: time="2024-09-16T15:31:5>
Sep 16 15:31:58 ip-172-31-27-176 containerd[4763]: time="2024-09-16T15:31:5>
Sep 16 15:31:58 ip-172-31-27-176 containerd[4763]: time="2024-09-16T15:31:5>
Sep 16 15:31:58 ip-172-31-27-176 containerd[4763]: time="2024-09-16T15:31:5>
Sep 16 15:31:58 ip-172-31-27-176 containerd[4763]: time="2024-09-16T15:31:5>
Sep 16 15:31:58 ip-172-31-27-176 systemd[1]: Started containerd.service - c>
```

### sudo apt-get install -y socat

```
ubuntu@ip-172-31-27-176:~$ sudo apt-get install -y socat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer requir
ed:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras
  docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  socat
0 upgraded, 1 newly installed, 0 to remove and 133 not upgraded.
Need to get 374 kB of archives.
After this operation, 1649 kB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 socat
amd64 1.8.0.0-4build3 [374 kB]
Fetched 374 kB in 0s (11.2 MB/s)
Selecting previously unselected package socat.
(Reading database ... 68108 files and directories currently installed.)
Preparing to unpack .../socat_1.8.0.0-4build3_amd64.deb ...
Unpacking socat (1.8.0.0-4build3) ...
Setting up socat (1.8.0.0-4build3) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-27-176:~$
```

**Step 6: Initialize the Kubecluster .Now Perform this Command only for Master.****sudo kubeadm init --pod-network-cidr=10.244.0.0/16**

```

ubuntu@ip-172-31-27-176:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W0916 15:39:33.685919 5313 checks.go:846] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container runtime is inconsistent with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.10" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-27-176 kubernetess kubernetess.default kubernetess.default.svc kubernetess.default.svc.cluster.local] and IPs [10.96.0.1 172.31.27.176]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-27-176 localhost] and IPs [172.31.27.176 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-27-176 localhost] and IPs [172.31.27.176 127.0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key
[kubeconfig] Using kubeconfig folder "/etc/kubernetes"
[kubeconfig] Writing "admin.conf" kubeconfig file
[kubeconfig] Writing "super-admin.conf" kubeconfig file
[kubeconfig] Writing "kubelet.conf" kubeconfig file
[kubeconfig] Writing "controller-manager.conf" kubeconfig file
[kubeconfig] Writing "scheduler.conf" kubeconfig file
[etcd] Creating static Pod manifest for local etcd in "/etc/kubernetes/manifests"
[control-plane] Using manifest folder "/etc/kubernetes/manifests"
[control-plane] Creating static Pod manifest for "kube-apiserver"
[control-plane] Creating static Pod manifest for "kube-controller-manager"
[control-plane] Creating static Pod manifest for "kube-scheduler"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Starting the kubelet
[wait-control-plane] Waiting for the kubelet to boot up the control plane as static Pods from directory "/etc/kubernetes/manifests"
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.001059178s

```

```

[api-check] Waiting for a healthy API server. This can take up to 4m0s
[api-check] The API server is healthy after 6.500965245s
[upload-config] Storing the configuration used in ConfigMap "kubeadm-config" in the "kube-system" Namespace
[kubelet] Creating a ConfigMap "kubelet-config" in namespace kube-system with the configuration for the kubelets in the cluster
[upload-certs] Skipping phase. Please see --upload-certs
[mark-control-plane] Marking the node ip-172-31-27-176 as control-plane by adding the labels: [node-role.kubernetes.io/control-plane node.kubernetes.io/exclude-from-external-load-balancers]
[mark-control-plane] Marking the node ip-172-31-27-176 as control-plane by adding the taints [node-role.kubernetes.io/control-plane:NoSchedule]
[bootstrap-token] Using token: ttay2x.n0squeukjai8sgfg3
[bootstrap-token] Configuring bootstrap tokens, cluster-info ConfigMap, RBAC Roles
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to get nodes
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to post CSRs in order for nodes to get long term certificate credentials
[bootstrap-token] Configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

```

Alternatively, if you are the root user, you can run:

```

export KUBECONFIG=/etc/kubernetes/admin.conf

```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at: <https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```

kubeadm join 172.31.27.176:6443 --token ttay2x.n0squeukjai8sgfg3 \
--discovery-token-ca-cert-hash sha256:d6fc5fb7e984c83e2807780047fec6c4f2acfe9da9184ecc028d77157608fbb6
ubuntu@ip-172-31-27-176:~$

```

**Run this command on master and also copy and save the Join command from above.**

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
ubuntu@ip-172-31-27-176:~$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
ubuntu@ip-172-31-27-176:~$
```

**Step 7: Now Run the command `kubectl get nodes` to see the nodes before executing Join command on nodes.**

```
Every 2.0s: kubectl get nodes ip-172-31-27-176: Mon Sep 16 15:45:34 2024
```

NAME	STATUS	ROLES	AGE	VERSION
ip-172-31-27-176	NotReady	control-plane	5m38s	v1.31.1

**Step 8: Now Run the following command on Node 1 and Node 2 to Join to master.**

```
sudo kubeadm join 172.31.27.176:6443 --token ttay2x.n0squeukjai8sgfg3 \
--discovery-token-ca-cert-hash sha256:d6fc5fb7e984c83e2807780047fec6c4f2acfe9da9184ecc028d77157608fbb6
```

**Node 1:**

```
ubuntu@ip-172-31-28-117:~$ sudo kubeadm join 172.31.27.176:6443 --token ttay2x.n0squeukjai8sgfg3 \
--discovery-token-ca-cert-hash sha256:d6fc5fb7e984c83e2807780047fec6c4f2acfe9da9184ecc028d77157608fbb6

[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 501.396793ms
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

ubuntu@ip-172-31-28-117:~$
```

**Node 2:**

```
ubuntu@ip-172-31-18-135:~$ sudo kubeadm join 172.31.27.176:6443 --token ttay2x.n0squeukjai8sgfg3 \
--discovery-token-ca-cert-hash sha256:d6fc5fb7e984c83e2807780047fec6c4f2acfe9da9184ecc028d77157608fbb6

[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.001003808s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

ubuntu@ip-172-31-18-135:~$
```

**Step 9: Now Run the command `kubectl get nodes` to see the nodes after executing Join command on nodes.**

```
ubuntu@ip-172-31-27-176:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-172-31-18-135	NotReady	<none>	88s	v1.31.1
ip-172-31-27-176	NotReady	control-plane	10m	v1.31.1
ip-172-31-28-117	NotReady	<none>	2m58s	v1.31.1

**Step 10: Since Status is NotReady we have to add a network plugin. And also we have to give the name to the nodes.**

**`kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml`**

```
ubuntu@ip-172-31-27-176:~$ kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
poddisruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
daemonset.apps/calico-node created
deployment.apps/calico-kube-controllers created
```

**`sudo systemctl status kubelet`**

```
ubuntu@ip-172-31-27-176:~$ sudo systemctl status kubelet
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/usr/lib/systemd/system/kubelet.service; enabled; preset: enabled)
   Drop-In: /usr/lib/systemd/system/kubelet.service.d
            └─10-kubeadm.conf
   Active: active (running) since Mon 2024-09-16 15:40:01 UTC; 11min ago
     Docs: https://kubernetes.io/docs/
   Main PID: 5989 (kubelet)
    Tasks: 10 (limit: 4676)
  Memory: 32.6M (peak: 33.2M)
     CPU: 10.705s
   CGroup: /system.slice/kubelet.service
            └─5989 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --config=/var/

Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497458      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s>
Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497516      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s>
Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497569      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s>
Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497620      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s>
Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497669      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s>
Sep 16 15:51:29 ip-172-31-27-176 kubelet[5989]: I0916 15:51:29.497719      5989 reconciler_common.go:245] "operationExecutor.VerifyControllerAttachedVolume s>
Sep 16 15:51:31 ip-172-31-27-176 kubelet[5989]: E0916 15:51:31.605091      5989 kubelet.go:2902] "Container runtime network not ready" networkReady="NetworkR>
Sep 16 15:51:32 ip-172-31-27-176 kubelet[5989]: I0916 15:51:32.366237      5989 scope.go:117] "RemoveContainer" containerID="f44f06967c5b3e567e07841a7b4352ae>
Sep 16 15:51:36 ip-172-31-27-176 kubelet[5989]: E0916 15:51:36.606675      5989 kubelet.go:2902] "Container runtime network not ready" networkReady="NetworkR>
Sep 16 15:51:41 ip-172-31-27-176 kubelet[5989]: E0916 15:51:41.608404      5989 kubelet.go:2902] "Container runtime network not ready" networkReady="NetworkR>
```



Now Run command `kubectl get nodes -o wide` we can see Status is ready.

```
ubuntu@ip-172-31-27-176:~$ kubectl get nodes -o wide
NAME                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-18-135    Ready     <none>    6m19s v1.31.1   172.31.18.135 <none>         Ubuntu 24.04 LTS     6.8.0-1012-aws     containerd://1.7.12
ip-172-31-27-176    Ready     control-plane 15m   v1.31.1   172.31.27.176 <none>         Ubuntu 24.04 LTS     6.8.0-1012-aws     containerd://1.7.12
ip-172-31-28-117    Ready     <none>    7m49s v1.31.1   172.31.28.117 <none>         Ubuntu 24.04 LTS     6.8.0-1012-aws     containerd://1.7.12
```

Now to Rename run this command

`kubectl label node ip-172-31-18-135 kubernetes.io/role=worker`

**Rename to Node 1:** `kubectl label node ip-172-31-28-117 kubernetes.io/role=Node1`

**Rename to Node 2:** `kubectl label node ip-172-31-18-135 kubernetes.io/role=Node2`

```
ubuntu@ip-172-31-27-176:~$ kubectl label node ip-172-31-28-117 kubernetes.io/role=Node1
node/ip-172-31-28-117 labeled
ubuntu@ip-172-31-27-176:~$ kubectl label node ip-172-31-18-135 kubernetes.io/role=Node2
node/ip-172-31-18-135 labeled
```

**Step 11: Run command `kubectl get nodes -o wide`. And Hence we can see we have Successfully connected Node 1 and Node 2 to the Master.**

```
ubuntu@ip-172-31-27-176:~$ kubectl get nodes -o wide
NAME                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-18-135    Ready     Node2     12m   v1.31.1   172.31.18.135 <none>         Ubuntu 24.04 LTS     6.8.0-1012-aws     containerd://1.7.12
ip-172-31-27-176    Ready     control-plane 21m   v1.31.1   172.31.27.176 <none>         Ubuntu 24.04 LTS     6.8.0-1012-aws     containerd://1.7.12
ip-172-31-28-117    Ready     Node1     13m   v1.31.1   172.31.28.117 <none>         Ubuntu 24.04 LTS     6.8.0-1012-aws     containerd://1.7.12
ubuntu@ip-172-31-27-176:~$
```

Or run `kubectl get nodes`

```
ubuntu@ip-172-31-27-176:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
ip-172-31-18-135    Ready     Node2     24m   v1.31.1
ip-172-31-27-176    Ready     control-plane 33m   v1.31.1
ip-172-31-28-117    Ready     Node1     25m   v1.31.1
ubuntu@ip-172-31-27-176:~$
```

**Conclusion:** In this experiment, we successfully set up a Kubernetes cluster with one master and two worker nodes on AWS EC2 instances. After installing Docker, Kubernetes tools (kubelet, kubeadm, kubectl), and containerd on all nodes, the master node was initialized and the worker nodes were joined to the cluster. Initially, the nodes were in the NotReady state, which was resolved by installing the Calico network plugin. We also labeled the nodes with appropriate roles (control-plane and worker). The cluster became fully functional with all nodes in the Ready state, demonstrating the successful configuration and orchestration of Kubernetes.